

Associate

The AI workspace built from scratch for lawyers. Not legal AI bolted onto Word. Not a chatbot pretending to know the law. The complete workspace where every kind of lawyer does their work — research, drafting, review, organisation — with an AI agent that holds context for every file in every matter, 24/7.

This document is the complete specification of what we are building, in the order it should be read.

Part One — The Idea

A senior lawyer's morning

A senior advocate in Bengaluru wakes up at 6:30 AM. By 7:00 she has finished her coffee. By 7:15 she is reading. Not the news — her brief. A senior associate has prepared an overnight memo: which matters need her attention today, what happened in the courts yesterday that affects her cases, which counterparty replied to which redline, where her draft submissions stand, what the cause-list looks like for the day. She reads it like a thoughtful piece of writing because it was written by someone who knows her, knows her cases, and respects her time.

This is how senior partners and GCs have always worked. They have associates who read everything, organise everything, draft the first version of everything. The partner does the judgment, the strategy, the relationships, the appearances — the irreducibly human work.

Associate gives every lawyer in India that experience. Not a chatbot. Not a search tool. Not another dashboard. An associate that reads, organises, drafts, watches, briefs — woven into a workspace that feels like the most thoughtfully designed legal publication you have ever read.

The single conviction

Existing legal AI is reactive. You ask, it answers. You draft, it suggests. You search, it returns. The model is fundamentally wrong because it treats the AI as a tool you reach for, not as the workspace you live in.

A workspace works for matters, not tasks. A lawyer doesn't think in tasks ("review contract", "research case"), they think in **matters** — the deal they're closing, the case they're litigating, the dispute they're arbitrating. A matter contains hundreds of files, dozens of communications, weeks of work, sometimes years of history. No existing tool holds the whole matter. They each see one slice and forget the rest by the next session.

Associate is what happens when we stop building tools and start building the workspace.

Three principles

Editorial over operational. When a lawyer opens Associate in the morning, they should feel like they are reading a thoughtful document prepared for them — not logging into a SaaS app. The

metaphor is *Financial Times Weekend*, not Salesforce. We present; we do not dashboard. There are no metric tiles. There are no KPI cards. Information lives in prose.

The agent is the author. Associate writes *into* the workspace. The matter brief, the footnotes, the overnight digest — the agent drafts these and the lawyer curates them. Our UI surfaces are places where the agent has already done work and the lawyer responds.

Words before widgets. A short paragraph that says "*counterparty sent revised redlines three days ago — the indemnity has been stripped of its cap*" is worth more than any status widget could be. Restraint is the aesthetic.

What Associate is, in one paragraph

Every matter — a case, a deal, a dispute, a transaction — becomes a **Living Brief**: a continuously-updated narrative document the agent maintains, where files appear inline as they become relevant, where the chronology writes itself, where new evidence updates the story automatically, where a floating chat is always present, and where the agent has read everything overnight so that you wake up already caught up. Documents live as full-fledged drafts in a beautiful editor with auto-citations and footnotes. Research happens in the same workspace, with answers grounded in real Indian case law and statutes. Translation runs in the background between English and India's major languages. The agent is everywhere. The lawyer never logs in; they walk in.

What Associate is *not*, by design

- **Not a dashboard.** No widgets, no charts, no overview screens. Information surfaces inside the work, where it is relevant.
- **Not a chatbot.** Conversation is the input mechanism for the agent, not the product surface. The product surface is the workspace.
- **Not a Word plugin.** We are the canvas. Word integration is at most an export option.
- **Not a CLM.** Contract lifecycles, signature workflows, vendor renewals live elsewhere. We integrate where we must.
- **Not a court-filing tool.** We do the desk work — research, drafting, review, organisation — that fills 70% of a lawyer's time. The lawyer files.
- **Not generic AI in a wrapper.** Every architectural decision was made for legal work specifically.

By saying no to these things, we say yes to the one thing that matters: being the workspace where every kind of lawyer does the work that AI can finally do for them.

Part Two — How Different Lawyers Use Associate

This is the most important part of the document. Associate is not a one-size-fits-all product layered on top of generic AI. It is **custom-shaped per lawyer**, because what an in-house counsel

needs is fundamentally different from what a litigator needs, which is fundamentally different from what a tax lawyer needs.

The same workspace, the same agent, the same context engine — but the matter templates, the default playbook, the surfaced information, the connectors, the document types, even the morning brief — all differ by archetype.

How customisation actually works

When a lawyer signs up, they answer one question: *"What kind of law do you practise?"* with structured options — In-house, Mid-tier firm associate, Solo practitioner, Senior partner, or Specialist (Arbitration / Tax / IP / Criminal / Family / Real Estate / M&A / Regulatory). They can pick one or more.

Based on this choice, Associate ships with:

- **A different starter playbook** — pre-loaded clauses, positions, templates relevant to that practice
- **Different default matter types** — a litigator gets criminal/civil/arbitration matters; an in-house counsel gets contract review/compliance/board matter types
- **Different connectors enabled by default** — eCourts integration is critical for litigators, not for transactional lawyers; Gmail and Drive matter for everyone
- **Different morning brief composition** — a litigator's "three things overnight" is heavy on judgments and cause-list; an in-house counsel's is heavy on regulator updates and counterparty replies
- **Different drafting templates** — pleading templates for litigators; MSA/SHA templates for transactional lawyers; reply-to-assessment templates for tax lawyers
- **Different research scope defaults** — criminal lawyers default to BNS/BNSS/BSA + state HC; tax lawyers default to ITAT + AAR + HC tax decisions

The aesthetic, the design language, the editorial voice never change. Every lawyer gets the same beautiful Associate. Only the substance underneath shifts to match how they work.

The eight day-in-the-life flows

Each is the way a real lawyer of that archetype actually spends their day, transformed by Associate. Read these end-to-end. They are the soul of the product.

Flow 1: The In-house Counsel at a startup

Today, without Associate:

- 9 AM: Open Outlook to 47 emails. CFO has forwarded 3 vendor MSAs marked "urgent". CEO wants a one-pager on whether they can cancel the AWS contract. Compliance team needs DPDP-compliance sign-off on a new SaaS vendor. HR wants 4 employment letters drafted for new hires.

- 12 PM: Still on the first MSA, copy-pasting from last quarter's deal because there's no playbook anywhere.
- 2 PM: Lunch at the desk while writing the AWS analysis.
- 4 PM: Realises the SEBI circular published this morning affects an ongoing fundraising discussion. Has to read the circular and re-write the response email.
- 8 PM: Still drafting employment letters. Goes home tired and behind.

With Associate:

- 6:30 AM: Phone notification from Associate's mobile app: "*Today's brief is ready.*" Reads it on the metro: 5 active matters need attention; the most urgent is the Pranos MSA where the agent has already drafted a DPDP-compliant counter-clause overnight; the AWS query has a draft response ready citing the relevant termination clause; the new SEBI circular has been pre-analysed and its implications written into the relevant matter's "Where we stand".
- 9 AM: Opens Associate at desk. Reviews and confirms the 4 approval cards on the Pranos draft (parties, dates, indemnity position, jurisdiction). The DPDP counter-clause is locked in.
- 9:30 AM: Floating chat: "*Draft 4 employment letters for these new hires.*" Drops in a CSV. Gets back 4 customised letters in 6 minutes, all referencing the company's standard employment terms playbook.
- 11 AM: Forwards the AWS response to CEO. Spends actual time on the strategic question of whether to cancel — not on the legal mechanics.
- 1 PM: Lunch. Real lunch. Outside.
- 4 PM: Pranos call. Has the negotiation brief on screen. Wins on the indemnity cap.
- 6 PM: Closes laptop. Sleeps well.

The shift: 3 hours of pure mechanical work disappear. The lawyer's day becomes *thinking* and *deciding*, not *typing*.

Flow 2: The Mid-tier Firm Associate

Today: drafts SHAs from term sheets manually, with a senior partner's checklist; researches case law in Manupatra/SCC Online with another tab open; tracks 12 active matters in a spreadsheet; partner sends edits at 11 PM that require redoing two clauses.

With Associate:

- Term sheet drops in via email forward → matter created automatically with parties, deal value, governing law, key dates parsed.
- "*Draft the SHA from this term sheet*" in chat → first draft in 6 minutes, structured per the firm's playbook (the firm playbook is loaded centrally for all associates), with all the standard clauses pre-filled, with 5 approval cards highlighting positions that need partner-level judgment.

- Associate reviews drafts, sends to partner with the 5 approval cards highlighted.
- Partner's edits come back; track-changes are inline, agent automatically updates the firm playbook with the new accepted positions.
- Research for the deal happens in the same workspace; cited judgments flow into footnotes automatically.
- 12 matters managed through 12 Living Briefs; partner asks "where are we on Sharma DD?"
→ `/status sharma` returns a 1-page brief in 30 seconds.

The shift: the associate produces work that looks like it came from a senior associate. Their effective output triples. Their billable leverage doubles. The path to senior associate gets faster.

Flow 3: The Solo Practitioner

Today: running a one-person practice. Drafting, research, court appearances, client management, billing — all alone. Cannot afford a junior associate (₹70,000/month minimum in metro India). Operating at ~40% of capacity because of admin overhead.

With Associate:

- Pays ₹2,500/month. Has an AI junior 24/7.
- Drafts 12 contracts a week instead of 4. Drafts pleadings, demand notices, replies in minutes.
- Walks into court with the matter brief on phone. The agent has prepared chronology, exhibits list, opposing counsel's prior arguments, judge's relevant past rulings.
- Client emails forwarded to the matter address get auto-categorised into the correct matter.
- Compliance calendar (when V1 of Compliance ships) keeps every client's filings on track.

The shift: operates with the leverage of a small team. Effective output triples. Stops turning down work because of capacity. This is the lawyer who most appreciates Associate viscerally — and the largest underserved segment of the Indian legal market.

Flow 4: The Litigation Specialist

Today: 312 active matters across 4 courts. Manages cause-list checking through eCourts every morning. Reads 40-page counter-affidavits and writes replies in long hand. Builds chronologies on legal pads. Maintains witness statements in physical files. Cross-examination prep is by candlelight, two days before the hearing.

With Associate:

- 6:00 AM: Agent has checked the cause-lists for all 4 courts overnight. Cross-referenced with the 312 matters. Surfaced the 8 matters listed today. Built the daily cause-list brief — what each matter is, what the prior orders said, what arguments are being heard.

- 7:00 AM: New counter-affidavit landed last night. Agent has read it (34 pages). Surfaces 2 new arguments and 1 internal contradiction with the witness's prior affidavit. *"Want me to draft the rejoinder?"*
- 8:30 AM: Agent has prepared cross-examination questions for today's PW-2 hearing — 47 questions cross-referenced against the witness's prior statements with exhibit pointers.
- In court: the matter brief is on tablet/phone. When the witness contradicts a prior statement, the lawyer taps and pulls up the exact prior statement with paragraph reference.
- Post-hearing: voice memo of the day's proceedings. Agent transcribes, files into the matter, updates the timeline.
- Evening: drafts closing submissions for a different matter using the chronology and evidence the agent has maintained. Citations auto-inserted.

The shift: the work that took weeks now takes days. The work that took days takes hours. The lawyer can take on more matters or do better work on the same matters. Most senior litigators tell us they would pay 10x what we charge for the cross-examination prep alone.

Flow 5: The Arbitration Specialist

Today: institutional arbitrations (SIAC, ICC, LCIA-India). 3,000-page hearing transcripts. Witness contradiction analysis. Award-writing for arbitral tribunals (when sitting as arbitrator).

With Associate:

- Drop in 3,000 pages of transcript: indexed in 15 minutes, fully searchable, with auto-built chronology.
- Find every contradiction across 4 witnesses' testimony in seconds.
- Generate dramatis personae automatically (parties, witnesses, counsel, experts, tribunal members).
- Closing submissions drafted with paragraph-level citations to specific evidence.
- For arbitrators: award skeleton drafted from the parties' submissions, leaving the reasoning sections for the arbitrator's own work.

Flow 6: The M&A / Due Diligence Lawyer

Today: a 400-document data room. 50-100 page DD reports. 4 weeks per deal at peak intensity.

With Associate:

- Index a 400-document data room in 15 minutes. Auto-categorise: corporate, contracts, employment, litigation, regulatory, IP, real estate, financials.
- Auto-generate the DD report skeleton in structured form with citations to specific data room documents.
- Cross-deal pattern recognition: *"This issue (related-party transactions documented informally) came up in 3 prior DDs you've done — here's how they ended."*

- Track deal-specific issues across multiple workstreams.

The shift: the DD that took 3 weeks takes 1 week. Quality goes up because the agent never gets tired and never misses a document.

Flow 7: The Tax Lawyer

Today: reads assessment orders, drafts replies in faceless assessments, researches ITAT/AAR/HC tax positions, tracks Finance Act amendments year-on-year.

With Associate:

- Faceless assessment notice received: parsed, parties identified, sections invoked, factual basis extracted.
- Reply drafted in 20 minutes — grounded in matter facts and prior orders, citing relevant ITAT/HC cases.
- Tax research with the full ITAT/AAR/CESTAT/HC corpus, paragraph-level citations.
- Finance Act amendments tracked daily; every active matter that touches an amended provision gets flagged.
- Interest computations, demand calculations, limitation tracking — automated.

Flow 8: The Criminal Lawyer

Today: reads FIRs, charge sheets, witness statements; drafts bail applications and replies; tracks hearing dates across multiple courts manually; prepares cross-examination through hand-marked transcripts.

With Associate:

- FIR ingested: parties, sections (BNS now, IPC for transitional matters), factual narrative, named witnesses auto-extracted.
- Bail application drafted in 15 minutes citing relevant SC/HC precedents on bail jurisprudence.
- Hearing dates pulled from eCourts automatically. Cause-list alerts the night before.
- BNS/BNSS/BSA cited correctly (and IPC/CrPC for ongoing matters that pre-date 1 July 2024).
- Cross-examination prep with witness contradictions surfaced from prior statements.
- Mobile-first because criminal lawyers live at courts, not desks.

What changes for everyone

Across every archetype, the pattern is consistent. Three hours per day come back. Quality goes up because the agent never gets tired. Mistakes go down because every claim is sourced. The

lawyer's mental energy moves from typing-and-formatting to thinking-and-deciding. This is the value proposition that makes lawyers — including the most sceptical senior counsel — pay willingly. Not "AI tool". "My associate."

Part Three — Where the Data Comes From

We can promise everything in Part Two only if we have the data infrastructure to back it up. This section is the complete enumeration of the Indian legal information sources we will integrate, in order.

The headline finding

Better than expected: most of what we need is already accessible. Three sources are particularly transformative:

eCourtsIndia.com API — a third-party enterprise-grade API on top of eCourts data. 26.7 crore+ case records across 600+ district courts, all 25 High Courts, the Supreme Court, NCLT, CAT, Consumer Forums, and tribunals. Per-request pricing, ₹200 free credits to start. **Has a working MCP server with 22 built-in tools** that connects directly to Claude. This is what we would have spent 12 months building in-house. We integrate from day one.

Indian Kanoon API — 1.4 million+ legal documents: every Supreme Court judgment from 1947, all 24 High Courts, 17 tribunals, Constitution, statutes, law commission reports. Documented public API at `api.indiankanoon.org`. Token-based, paid per query. This is our foundational legal corpus for grounding every research answer.

Open Government Data Platform (data.gov.in) — 100,000+ government datasets across all ministries. Free with API key registration. Has Python and R libraries. Backbone of our regulator radar.

Tier 1 — Integrate during the demo (months 0-3)

Source	What it covers	Access
Indian Kanoon	1.4M+ judgments, statutes, Constitution	https://api.indiankanoon.org/documentation/
eCourtsIndia API	26.7Cr+ records across all courts + tribunals	https://ecourtsindia.com/api
OGD Platform	100,000+ govt datasets	https://www.data.gov.in/apis
India Code	All Central + State + UT acts + subordinate legislation	https://www.indiacode.nic.in/ (scrape — 4 weeks dev)
Supreme Court of India	Daily orders, judgments, cause lists, transcripts	https://www.sci.gov.in/ (eCourts API covers most)

Tier 2 — Integrate during MVP (months 3-9)

Regulator websites without APIs. Build a normalised scraper layer.

Financial regulators:

- SEBI — circulars, master circulars, all notifications: <https://www.sebi.gov.in/>
- RBI — circular index, master directions: <https://www.rbi.org.in/>
- IRDAI: <https://www.irdai.gov.in/>
- PFRDA: <https://www.pfrda.org.in/>
- CCI: <https://www.cci.gov.in/>

Tax authorities:

- CBDT (Income Tax): <https://incometaxindia.gov.in/Pages/communications.aspx>
- CBIC (GST/Customs/Excise): <https://www.cbic.gov.in/>
- GST Portal: <https://www.gst.gov.in/> — also licensed access via GST Suvidha Providers

Corporate & insolvency:

- MCA: <https://www.mca.gov.in/> — no official API, integrate via licensed third-parties (AuthBridge, Surepass, Technowire) at ~₹2 per CIN lookup
- IBBI: <https://ibbi.gov.in/>

Other ministries:

- MeitY (DPDP, IT Act): <https://www.meity.gov.in/>
- DGFT (Foreign Trade): <https://dgft.gov.in/>
- Ministry of Labour & Employment: <https://labour.gov.in/>
- Ministry of Environment, Forest, Climate: <https://moef.gov.in/>
- TRAI (Telecom): <https://www.trai.gov.in/>
- FSSAI (Food): <https://www.fssai.gov.in/>
- BIS (Standards): <https://www.bis.gov.in/>

Tier 3 — Tribunals (during V1, mostly via eCourts API)

NCLT, NCLAT, ITAT, CESTAT, SAT, DRT/DRAT, TDSAT, AFT, NCDRC, NGT — every tribunal that matters. eCourts API covers most. Direct site integration as fallback.

Tier 4 — Editorial sources (RSS, ongoing)

LiveLaw, Bar and Bench, TaxGuru, IBC Laws — for editorial framing in the morning brief. RSS available; cheap to integrate.

Tier 5 — Procedural rules (V1)

SC Rules 2013, every HC's procedural rules, CPC, CrPC, BNSS, BNS, BSA, Limitation Act, Arbitration Act, NI Act — via India Code scraper.

Tier 6 — Bar & institutional (V2, foundation for Lawyer Graph)

Bar Council of India for verification, NALSA for legal aid integrations, Indian Law Institute for content.

Tier 7 — Subscription databases (selective, post-Series A)

SCC Online, Manupatra, Lexis India, Westlaw India, Taxmann. **Strategy: bring-your-own-database.** Firms with existing subscriptions authenticate and we proxy queries through their session. We never license these ourselves; we ride on top of customer licences. This avoids a content-licensing war and makes us additive rather than competitive to incumbent databases.

Tier 8 — E-filing portals (V3)

SC AOR e-Filing, every HC e-filing portal, tribunal e-filing systems. We do not file in V1; we generate print-ready or e-filing-ready output.

Tier 9 — Adjacent infrastructure APIs

DigiLocker for document verification; Aadhaar/PAN/GST verification for client KYC; CIBIL/CRIF for recovery cases; GST Suvidha Provider APIs for licensed GSTN access; DSC verification APIs.

Indian-built where possible

For the AI layer itself, we partner with Indian infrastructure where it makes sense:

- **Sarvam, Krutrim, or self-hosted IndicTrans2** for Indian-language translation
- **AWS Mumbai region** for hosting (data residency requirement)
- **Indian payment rails** (Razorpay, Cashfree)
- **Indian KYC providers** (Surepass, AuthBridge, Cashfree Verification)

The narrative — *"your legal AI is Indian, and it knows that"* — is part of why senior counsel choose us. Sovereignty matters in legal.

Part Four — Guardrails and Necessities

This is the section where we admit what could go wrong and explain how we prevent it. Every legal AI tool that has failed publicly has failed on one of these axes. We engineer against each from day one.

Hallucination is a fire-able offence

The Bombay High Court has fined lawyers ₹50,000 for submitting AI-fabricated citations. A trial

court in Andhra Pradesh decided a property dispute citing four synthetic Supreme Court rulings that did not exist. The "Mercy vs Mankind" incident has become a cautionary tale across Indian bar associations. The message from courts is consistent: **if you submit it, you own it.**

This means hallucination is not a quality issue for Associate. It is an existential issue. A single high-profile incident can end the company.

How we prevent it:

1. **Architecture-level citation enforcement.** The agent cannot return a research answer or insert a citation in a draft without a verifiable source. If no source exists, the system says so and refuses to invent one.
2. **Every citation is clickable and verifiable.** The lawyer can tap any case name and see the actual judgment text from Indian Kanoon or eCourts. We never display a citation we have not retrieved from a real source.
3. **Provenance markers throughout.** Every paragraph in a generated document carries a small inline indicator: agent-drafted, lawyer-written, lawyer-edited-from-agent. Audit trail is built into the data model, not bolted on.
4. **Confidence thresholds.** When the agent's certainty is below threshold, it does not bluff. It surfaces the uncertainty and asks. *"I couldn't find a direct precedent on this point. Want me to search more broadly, or do you want to reason from first principles?"*
5. **Lawyer-responsibility framing.** Every draft includes a footer: *"AI-assisted. Final responsibility rests with the signing lawyer."* This is not a disclaimer for our protection; it is a reminder of the standard of care that legal practice demands.

Data residency and confidentiality

Legal work is privileged. Our data architecture must not just respect privilege — it must be defensibly architected to demonstrate respect.

Hard rules:

- All data in **AWS Mumbai region**, never replicated outside India without explicit consent.
- Per-matter encryption keys. No cross-matter data flow without explicit linkage by the lawyer.
- **SOC 2 + ISO 27001** by Month 12 of the public product.
- **Optional on-premise / private VPC deployment** for highly sensitive matters using local quantised models. We commit to this for enterprise customers.
- We do not train our models on customer data without explicit opt-in. Default is opt-out.
- We do not sell, share, or analyse customer data for advertising or any non-product purpose. Ever.

Bar Council compliance

The Bar Council of India has been actively litigating against legal-tech tools that aggregate lawyer data, recommend specific lawyers to clients, or run touting-adjacent operations. Some platforms have been shut down.

How we comply:

- We engage with the BCI proactively from day one. We position the verified-profile system as helping the BCI track its registered members and enforce its rules.
- We do not allow lawyer-to-client direct touting through the Lawyer Graph.
- We do not surface judges' rulings as predictive data points for individual cases. Aggregate-only, framed as research support.
- We do not publish ranking, comparison, or "best lawyer" lists.
- We respect every takedown request from the BCI or any regulator, no questions asked.

Court analytics — the line we do not cross

It is possible to build analytics that profile individual judges' rulings ("Justice X has ruled against bail in 73% of cases involving public servants"). We will not.

Why: this kind of analysis is on the edge of contempt of court in India. It also could be misused — by parties seeking to pressure judges, by media seeking sensational angles, by political actors with bad intent. We do not want to be the platform that enabled that.

What we do instead: aggregate signals at the court-level or bench-level (not individual judge level), framed as research support, never as prediction. *"Bombay HC commercial division typically takes 14-18 months on contract disputes"* is acceptable. *"Justice X is favourable to claimants"* is not.

Hard rules around child safety and harmful content

Even though our customers are professional lawyers, we ship a product accessible to anyone. Our content systems strictly refuse to produce:

- Content that sexualises or grooms minors, in any framing, ever.
- Content that facilitates concrete harm — weapons, violence, criminal facilitation.
- Content that is defamatory of identifiable people without legal grounding.
- Content that could be construed as legal advice to non-lawyers (we serve lawyers; we do not directly advise pro-se litigants).

These rules apply at the model level. They are not optional features.

Professional liability

Every legal product carries malpractice risk. Our position: we are a tool that augments a lawyer's work, never a replacement for legal judgment. The lawyer is always the responsible party. This is reinforced in:

- Our terms of service
- Our UI (the "Final responsibility rests with the signing lawyer" footer)
- Our marketing (we never claim Associate "practises law" or "gives legal advice")
- Our pricing (we charge lawyers, not their clients)

Insurance partnership (Phase 4): once at scale, we partner with PI insurance carriers. The audit trail Associate produces *reduces* a lawyer's liability exposure because every action is timestamped, every citation is sourced, every draft has provenance. Insurers should price our customers favourably — and we should make this contractual.

Failure modes we plan for

Risk	Likelihood	Mitigation
eCourts portal structure changes	Medium	Use third-party API as primary; maintain backup scraper; engineering on call
Bar Council restricts platform	Low– Medium	Proactive engagement; BCI co-design; rapid takedown compliance
Hallucination incident in court	Low	Architecture-level prevention; lawyer-responsibility framing; insurance backstop
Regulator (DPDP authority) issues restrictive guidelines	Medium	Be the most compliant tool in the market; engage proactively; lawyer-PM advises
SCC Online / Manupatra build competing AI	High	BYO database strategy; we ride their licences instead of competing
Harvey enters India aggressively	High	Unit economics work at ₹2,500/lawyer; their cost structure cannot match
Over-build the demo, run out of runway	High	Hard time-box demo to 12 weeks; ship with imperfections; polish in MVP

We discuss these openly with investors. They appreciate honesty here far more than confidence.

Part Five — Confidence Signals That Sell

When we walk into a room — investor meeting, design partner pitch, customer demo, conference panel — these are the signals we use to convey conviction. Every signal is a real artefact, not a marketing claim.

The 60-second magic moment

In any pitch — investor or customer — we do this within the first 60 seconds:

1. Open Associate cold (no pre-loaded matter)
2. Drop in a real PDF (a real contract, a real FIR, a real court order)
3. Within 6 seconds, the matter is created with parties, dates, jurisdiction, regulatory surface, and a first "Where we stand" agent-written narrative
4. Ask a research question in the floating chat
5. Get back a real, cited Indian-law answer with paragraph-level citations to actual judgments

By second 60, the room has changed. Sceptics turn into believers. We have done this with senior partners, GCs, and investors. It works every time, because it is real.

The aesthetic itself

The product looks like nothing else in legal tech. No SaaS dashboard, no dark-mode-and-purple-accent vibes, no Material-Design-by-default. Editorial typography, warm paper-like surfaces, restrained colour palette, generous whitespace, prose-first information design.

The aesthetic itself is a confidence signal. Senior counsel respond viscerally to design quality. They have spent careers reading well-typeset judgments and elegantly-drafted contracts. When they see Associate, they see a product made by people who respect the work. They take us seriously immediately.

Real Indian sources, named explicitly

Every demo and pitch should name our sources by name:

- *"This research is grounded in Indian Kanoon's full corpus — 1.4 million judgments."*
- *"Court data comes live from eCourts India's API — 26.7 crore case records."*
- *"Regulatory updates flow from SEBI, RBI, MCA, CBDT — refreshed daily."*
- *"Translation runs on India-built models, not generic Google Translate."*

This is the *built-in-India-from-India* narrative made concrete. It is also a moat — every named source is one we have integrated, and integrating it is hard work. Competitors (especially foreign ones) have not done it.

The design partner cohort

By Month 6, we have 50 hand-picked design partners using Associate. Senior advocates, GCs at notable startups, partners at recognisable mid-tier firms. We collect quotes — real ones, with names, with their permission.

"Associate has given me three hours back every day." — a real GC quote. *"My associate has read everything before I have."* — a real arbitration partner quote.

We use these quotes. We also stay in close contact with each design partner because they are our most credible sales channel. A senior advocate who tells five other senior advocates about Associate is worth more than any LinkedIn post.

Integration receipts

When pitching, show actual screenshots and demo of:

- A real Indian Kanoon citation rendering inline in a draft
- A real eCourts cause-list pulled live for a real case number
- A real SEBI circular ingested overnight, written into a matter's brief
- A real Hindi-to-English bilingual rendering of a court order

These are not Photoshopped. They are running on the demo server. Investors and customers can ask to see any of them, any time, on any case number.

The unit economics conversation

Investors care about gross margin, retention, payback period. We have answers:

- **Gross margin:** ~75% (LLM costs are <15% of revenue at our pricing; infra <8%; data licensing <2%).
- **Retention:** in legal SaaS the benchmark is 90%+ NRR for strong products. We design for higher because of compounding playbook lock-in.
- **Payback period:** under 6 months at our planned ARPU and CAC, validated by design partner conversion rates.
- **TAM:** 1.5M lawyers in India $\times$ ₹3,000/month average ARPU = ₹540 Cr ARR ceiling on the seat-based business alone. Add Compliance Engine, Lawyer Graph monetisation, and platform layers, and the addressable opportunity is ₹3,000–5,000 Cr in India alone, with international meaningfully larger.

The team narrative

We are a team that has built a real product, in India, for India, with deep design quality and real engineering. Not a "GPT wrapper" startup. Not a foreign-tech-bolted-on-India startup. We are the team that did the hard work of integrating eCourts, scraping the regulators, learning the playbooks, building the editorial product.

This narrative is not vanity. It is signal. Investors and customers both pattern-match on team strength as the strongest predictor of outcome. We earn the pattern-match by doing the work.

The clarity of what we are not

When we say "*we are not a chatbot, not a dashboard, not a Word plugin, not a CLM, not a court-filing tool*" — and we mean it, and our product visibly reflects it — we communicate **conviction**.

Most products try to be everything. We are explicit about what we refuse to be. Sophisticated buyers and investors notice this.

The roadmap that compounds

We do not promise everything in V1. We promise a credible roadmap where each phase compounds the next. Layer 1 (the workspace) wins users. Layer 2 (the Lawyer Graph) locks them in. Layer 3 (Compliance) expands revenue. Layer 4 (platform expansion) makes us the OS. We tell this story clearly. People believe sequenced plans more than they believe big-bang plans.

The proof that we know our customer

Throughout every conversation, we use Indian legal terminology accurately. We know that a *vakalatnama* is not just a "power of attorney" — it is the specific authorisation a lawyer files. We know that a Section 138 NI Act notice is what cheque-bounce cases turn on. We know what *zamānat* means in court Hindi versus *jamānat* in colloquial Hindi. We know which HC has Original Side rules and which doesn't. We know the difference between SIAC and ICC arbitration timelines.

This terminological fluency is itself a confidence signal. It tells lawyers we are not generalists who happened to choose legal as a market. We are people who have done the work to understand their work.

Part Six — The Depth (How We Build)

The roadmap, phased

Phase 0 — Demo (months 0–3)

Goal: raise pre-seed/seed, sign 50 design partners, validate willingness to pay.

Build:

- Full workspace UI (Home, Matter Brief, Document Editor, Floating Chat, Matter Rail) with editorial design language
- Voice/text matter intake with parsed-matter generation
- Document upload + OCR + entity extraction
- Indian Kanoon API integration for grounded research with footnote citations
- eCourtsIndia API integration via MCP (cached for ~50 demo cases)
- OGD + targeted scraped regulator data for "overnight three things"
- Coming-soon previews for Compliance, Lawyer Graph, real-time eCourts, full Translation
- Hindi-only translation for two demo documents (proof-of-capability)

Don't build: real billing, real client portal, real e-filing, anything beyond Hindi translation, real Lawyer Graph at scale.

Phase 1 — Closed Beta MVP (months 3-9)

Goal: convert 30 of the 50 design partners to paying users at ₹2,500/month. Validate retention.

Build:

- Full eCourts API integration across all 600+ courts (live, not cached)
- Indian Kanoon paid tier with full search and retrieval
- Comprehensive India Code scraper (every Central Act + amendments)
- All Tier 2 regulator scrapers (SEBI, RBI, MCA via licensed third-party, IRDAI, PFRDA, CCI, CBDT, CBIC, MeitY, DGFT)
- All Tier 3 tribunals via eCourts API
- Lawyer Graph V1: verified profiles, co-counsel collaboration on shared matters
- Translation V1: English ↔ Hindi document-level + side-by-side bilingual
- Compliance Engine: still demo-only, but waitlist closed and design partners shown deeper previews

Team: 5 engineers + 1 designer + 1 product + 2 GTM. Burn ~₹40L/month. Seed-funded.

Phase 2 — V1 Public Launch (months 9-18)

Goal: scale from 30 paying users to 1,000. Establish category leadership.

Build:

- Compliance Engine V1 (real, not demo): entity profiles, calendar generation, auto-drafted forms (no submission yet — draft + lawyer-files-manually)
- Lawyer Graph V2: full multi-party matters including opposing counsel, counterparty memory aggregation, firm pages
- Translation V2: add Tamil, Bengali, Marathi. Add audio transcription with translation. Witness statement workflow.
- Bring-your-own SCC Online and Manupatra connectors
- Mobile app (iOS + Android)

Team: 12 engineers + 2 designers + 2 product + 5 GTM + 2 ops. Burn ~₹1.5 Cr/month. Series A.

Phase 3 — Platform Expansion (months 18-30)

Goal: 25,000+ Indian lawyers. Begin selective international expansion (UAE, Singapore Indian-diaspora work).

Build:

- Compliance Engine V2: full submission flows where APIs exist (CBDT, GSTN). Chrome extension for portal-based filings. SEBI LODR for listed.
- Lawyer Graph V3: bench analytics (carefully framed)
- Translation V3: real-time courtroom translation. Mobile-first mode. Certified translation tier.
- Full e-filing-prep coverage
- Client portal for in-house teams
- Billing/time tracking
- API for third-party integrations

Team: 30+ engineers, 5+ designers, full GTM org. Series B.

Phase 4 — The OS (months 30+)

Insurance layer, legal recruitment, CLE, marketplace, broader API ecosystem. The decade play.

The team we need

In priority of difficulty to hire:

1. **Search/IR engineer who understands legal text.** Indian Kanoon's founder, or someone like him.
2. **Data ingestion engineer** who can scrape government portals at scale without breaking.
3. **ML engineer** focused on legal-aware chunking, citation enforcement, hallucination prevention.
4. **Legal product manager** who is an actual practising lawyer or recently-was.
5. **Founding designer** who understands editorial design, not generic SaaS. We have this; protect it.

The architecture

Five layers, each independently strong, together unprecedented in Indian legal tech:

Layer 1 — Context Engine. Persistent matter memory. Vector-embedded with legal-aware chunking. Built once, used everywhere.

Layer 2 — Indian Legal Corpus. India Code + 1.4M+ Indian Kanoon judgments + all major regulations. Vector-embedded, refreshed daily.

Layer 3 — The Agent. LLM-orchestrated reasoning over matter context, library, and playbook simultaneously. Citation enforcement at architecture level.

Layer 4 — The Workspace. Web-native canvas built on TipTap. Native rich text, track changes, comments, version history, Cursor-style inline AI suggestions.

Layer 5 — Data Privacy. AWS Mumbai region. Per-matter encryption keys. SOC 2 + ISO 27001 by Month 12. Optional on-premise / private VPC for sensitive matters using local quantised models.

The single-paragraph promise

If everything in this document goes right, in 36 months Associate is what every law student in India is trained on, what every solo lawyer opens in the morning, what every firm runs on, and what every in-house counsel uses to manage compliance. The data infrastructure described in Part Three is the moat. The editorial workspace experience described in Parts One and Two is the wedge. The guardrails in Part Four make us trustable. The signals in Part Five make us pitchable. The depth in Part Six makes us buildable.

We are not building an AI tool. We are building **the operating system of Indian legal practice** — layer by layer, source by source, lawyer by lawyer.

Live document. Updated as we learn. Owners: founders. When in doubt: the workspace philosophy in Part One is correct, and the demo in Phase 0 is the proof.